

CS344 : Operating systems Lab

Assignment-1

Group Details

Member - 1

Name: Dhruvesh Bhure

Roll No: 200123018

Member - 2

Name: Soham Roy

Roll No: 200123055

Member - 3

Name: Dev Sandip Shah

Roll No: 200123074

Part 1: Kernel threads

We are asked to implement 3 new system calls to simulate the working of threads.

1. ***thread_create()*** - function to create a kernel thread.
2. ***thread_join()*** - function to wait for the thread to finish.
3. ***thread_exit()*** - function that allows the thread to exit.

Structure of a process -

```
37 // Per-process state
38 struct proc {
39     uint sz; // Size of process memory (bytes)
40     pde_t* pgdir; // Page table
41     char *kstack; // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid; // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan; // If non-zero, sleeping on chan
48     int killed; // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd; // Current directory
51     char name[16];
52     int isThread; // Process name (debugging) // implemented by us denotes whether
    the process is thread or not
53 };
54
```

The variable ***isThread*** on the last line is added by us in the struct to indicate whether the current process is thread or not.

thread_create() system call looks like this:

int thread_create(void(*fcn)(void*), void *arg, void*stack)

Where *fcn* is the pointer to the custom function on which thread is created, *arg* is the arguments required during the function running, *stack* is manually created for thread execution.

This call creates a new kernel thread which shares the address space with the calling process.

```
93 int sys_thread_create(void){
94     void (*fcn)(void*), *arg, *stack;
95     argptr(0, (void*) &fcn, sizeof(void*)(void *));
96     argptr(1, (void*) &arg, sizeof(void*));
97     argptr(2, (void*) &stack, sizeof(void *));
98     return thread_create(fcn, arg, stack);
99 }
```

The code above is implemented in *sysproc.c*
thread_create() is a user call which is implemented by us in *proc.c* file which makes a system call to *sys_thread_create()* which takes the arguments from *thread_create()* and makes a system call by calling the *thread_create()* function from the kernel.

```

537
538 int thread_create(void (*fcn)(void *),void *arg,void* stack){
539     if((uint)stack==0) // if no memory is allocated to the stack return -1
540     {
541         return -1;
542     }
543     int i,pid; // pid is the process id
544     struct proc *newproc; //new process / child process
545     struct proc *current_proc = myproc(); //Current process in which thread is being created
546
547     if((newproc=allocproc())==0)return -1; // allocating a new process to child if not successful return -1
548     newproc->pgdir = current_proc->pgdir; // making sure both have same page directory
549     newproc->sz = current_proc->sz; // making sure both have same size
550     newproc->parent = current_proc; // pointing parent to child
551     *newproc->tf = *current_proc->tf; // same trap frame
552
553     newproc->isThread = 1; // telling new process it is a thread
554
555     newproc->tf->eax = 0; // when new process gets finished 0 will be returned
556
557     newproc->tf->eip = (int)fcn; // this is the function on which thread will work
558
559     newproc->tf->esp = (int) stack + 4096;
560     newproc->tf->esp -= 4;
561     *((int*)(newproc->tf->esp)) = (int) arg;
562     newproc->tf->esp+=4;
563     *((int*)(newproc->tf->esp)) = 0xffffffff;
564     for(i=0;i<NOFILE;++i)
565     {
566         if(current_proc->ofile[i])
567         {
568             newproc->ofile[i] = filedup(current_proc->ofile[i]); // copying all opened files from current
process to new process
569         }
570     }
571     newproc->cwd = idup(current_proc->cwd);
572     safestrcpy(newproc->name,current_proc->name,sizeof(current_proc->name));
573     pid = newproc->pid;
574     acquire(&ptable.lock);
575     newproc->state = RUNNABLE;
576     release(&ptable.lock);
577     return pid;
578 }

```

In the above code, current process_id(*pid*) is copied into a newly created thread, then page_directory(*pgdir*) and size(*sz*) are also copied. Making customised *isThread* variable 1 to denote it a thread.

In the trapframe(*tf*) of the new process(*newproc*) we have initialised *eax*, *eip*, *esp*. Then we iterated over all the files in the current process and copied those files that were open to the new process(*newproc*).

The other new system call is :

int thread_join(void)

This call waits for a child thread that shares the address space with the calling process. It returns the *pid* of the waited-for child or -1 if none.

```
101 int sys_thread_join(void){
102     return thread_join();
103 }
```

The code above is implemented in *sysproc.c*

thread_join() is a user call which is implemented by us in *proc.c* file which makes a system call to *sys_thread_join()*.

```
580 int thread_join(void){
581     struct proc *i;
582     int havekids; // havekids is the boolean value to check if the thread has any child or not
583     int pid;
584     struct proc *current_proc = myproc(); //Current process in which thread is being created
585
586     acquire(&ptable.lock);
587     while(1){
588         havekids = 0;
589         for(i = ptable.proc; i < &ptable.proc[NPROC]; ++i){
590             if(i->isThread != 0 && i->parent == current_proc){ //if the process is a thread and child of
the current process
591                 havekids = 1;
592                 if(i->state == ZOMBIE){ //if it's in zombie state then re-initialize to make it
available to be used by other processes and return
593                     pid = i->pid;
594                     i->kstack = 0;
595                     i->pid = 0;
596                     i->parent = 0;
597                     i->name[0] = 0;
598                     i->killed = 0;
599                     i->state = UNUSED;
600                     release(&ptable.lock);
601                     return pid;
602                 }
603             }
604         }
605         if(!havekids || current_proc->killed){ // if it doesn't have any child or the current process itself
is aborted then return -1
607             release(&ptable.lock);
608             return -1;
609         }
610         sleep(current_proc, &ptable.lock); // wait for any thread to complete its execution
611     }
612 }
613 }
```

The other new system call is :

int thread_exit(void)

This system call allows a thread to terminate.

```
105
106 int sys_thread_exit(void){
107     return thread_exit();
108 }
```

The code above is implemented in *sysproc.c*

thread_exit() is a user call which is implemented by us in *proc.c* file which makes a system call to *sys_thread_exit()*.

```
616 int thread_exit(){
617     struct proc *current_proc = myproc();
618     struct proc *i;
619     int file;
620
621     if(current_proc==initproc)
622     {
623         panic("exit init");
624     }
625     for(file = 0; file<NOFILE; file++){
626         if(current_proc->ofile[file]){
627             fileclose(current_proc->ofile[file]); // closing all opened files of current process
628             current_proc->ofile[file] = 0;
629         }
630     }
631     begin_op();
632     iput(current_proc->cwd);
633     end_op();
634     current_proc->cwd = 0;
635     acquire(&ptable.lock);
636
637     wakeup1(current_proc->parent);
638
639     for(i=ptable.proc; i<&ptable.proc[NPROC]; ++i){
640         if(i->parent==current_proc){
641             i->parent = initproc;
642             if(i->state==ZOMBIE){
643                 wakeup1(initproc);
644             }
645         }
646     }
647     current_proc->state = ZOMBIE;
648     sched();
649     panic("exit zombie");
650 }
```

```
QEMU
Machine View
Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ thread
Starting do_work: s:b2
Starting do_work: s:b1
Done s:2F9C
Done s:2F78
Threads finished: (4):5, (5):4, shared balance:3218
$ thread
StartStarting do_work: s:b2
ing do_work: s:b1
Done s:2F9C
Done s:2F78
Threads finished: (7):8, (8):7, shared balance:3207
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F9C
Done s:2F78
Threads finished: (10):11, (11):10, shared balance:3200
$ _
```

The above output is wrong (because $3200 + 2800$ is 6000 but the received outputs are different every time we run the `thread.c` file). The reason behind this is that the 2 threads are changing the value of the same global variable `total_balance`. It might happen that both threads read an old value of the `total_balance` at the same time, and then update it at almost the same time as well. As a result the deposit (the increment of the balance) from one of the *threads* is lost. This can be fixed using *synchronisation*.

Part 2: Synchronisation

To fix this synchronisation error we have implemented a ***spinlock*** that allows us to execute the update atomically.

Struct of *spinlock* :

```
5 struct thread_spinlock{
6     volatile uint lock;
7     char *name;
8 };
```

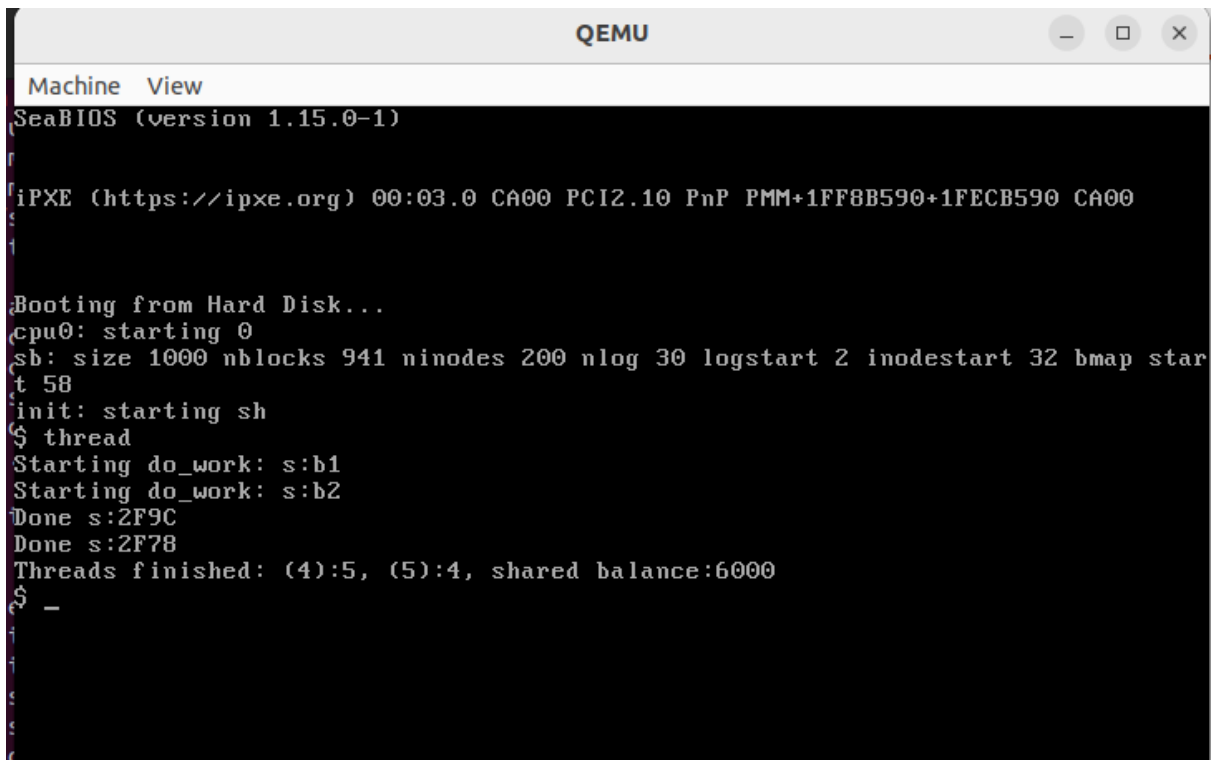
There are 3 functions implemented in *spinlock* -

1. ***void thread_spin_init(struct thread_spinlock *lk)*** - initialise the spinlock to the correct initial state.
2. ***void thread_spin_lock(struct thread_spinlock *lk)*** - a function to acquire a spinlock.
3. ***void thread_spin_unlock(struct thread_spinlock *lk)*** - a function to release the spinlock.

Implementation of these 3 functions are as follows :

```
25 void thread_spin_init(struct thread_spinlock *lk){
26     lk->lock = 0;
27     lk->name = "null";
28 }
29
30 void thread_spin_lock(struct thread_spinlock *lk){
31     while(xchg(&lk->lock,1)!=0);
32     __sync_synchronize();
33 }
34
35 void thread_spin_unlock(struct thread_spinlock *lk){
36     __sync_synchronize();
37     asm volatile("movl $0, %0" : "+m" (lk->lock) : );
38 }
```

Output using spinlock :



```
Machine View
SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B590+1FECB590 CA00

Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F9C
Done s:2F78
Threads finished: (4):5, (5):4, shared balance:6000
$ -
```

So using spinlock we got the correct output.

When the CPU count is 2, all threads of the process run in parallel on different CPUs, spinlocks are perfect, each process enters a short critical section, updates the shared balance atomically and then releases the spinlock for other threads to make progress.

```
220 ifndef CPUS
221 CPUS := 2
222 endif
```

However, when we change the number of CPU's in the Makefile from 2 to 1, then spinlock takes more time because all threads of the process start to spin endlessly, waiting for the interrupted (lock-holding) thread to be rescheduled and run again the spinlocks become inefficient.

```
220 ifndef CPUS
221 CPUS := 1
222 endif
```


To fix the above problem, we have implemented ***mutex lock***.

Struct of *mutex* :

```
10 struct thread_mutex{
11     volatile uint lock;
12     char *name;
13 };
```

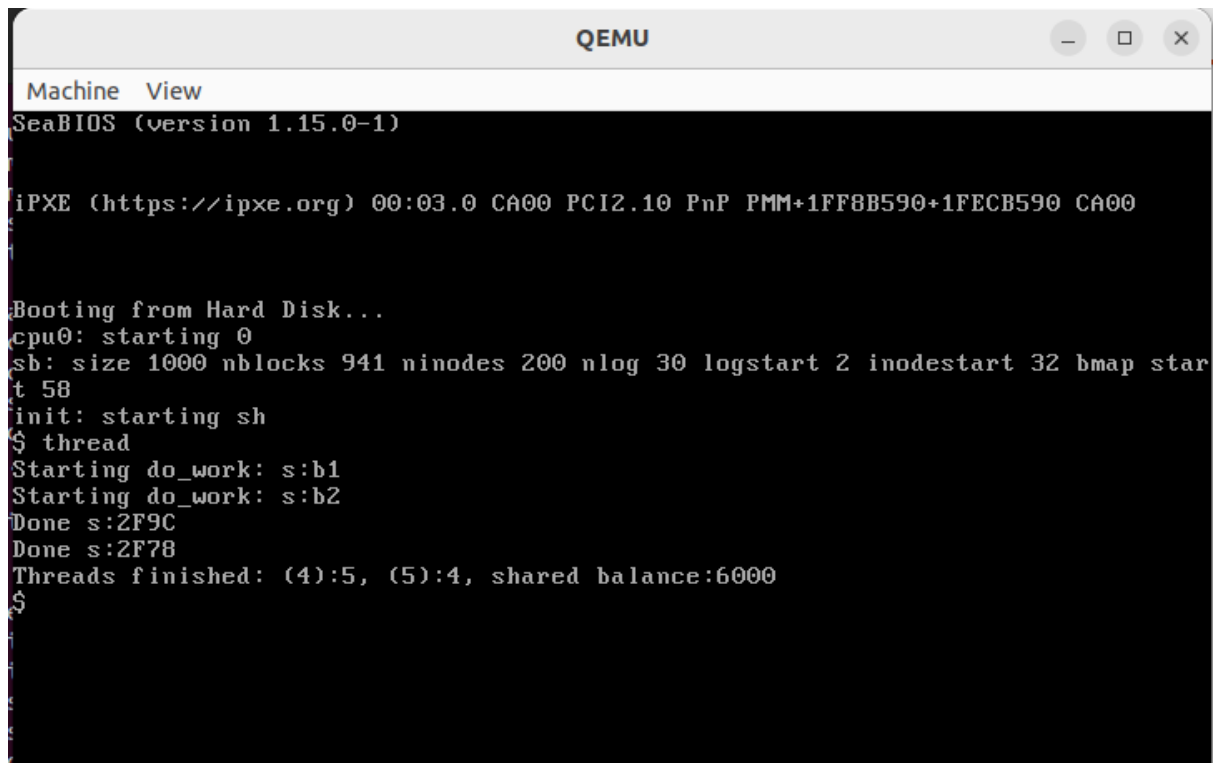
There are 3 functions implemented in *mutex* -

1. ***void thread_mutex_init(struct thread_mutex *m)*** - initialise the mutex lock to the correct initial state.
2. ***void thread_mutex_lock(struct thread_mutex *m)*** - a function to acquire a mutex lock.
3. ***void thread_mutex_unlock(struct thread_mutex *m)*** - a function to release the mutex lock.

Implementation of these 3 functions are as follows :

```
40 void thread_mutex_init(struct thread_mutex *lk){
41     lk->lock = 0;
42     lk->name = "null";
43 }
44
45 void thread_mutex_lock(struct thread_mutex *lk){
46     while(xchg(&lk->lock,1)!=0){
47         sleep(1);
48     }
49
50     __sync_synchronize();
51 }
52
53 void thread_mutex_unlock(struct thread_mutex *lk){
54     __sync_synchronize();
55     asm volatile("movl $0, %0" : "+m" (lk->lock) : );
56 }
```

Output using mutex :



```
Machine View
SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B590+1FECB590 CA00

Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F9C
Done s:2F78
Threads finished: (4):5, (5):4, shared balance:6000
$
```

Observations :

We observed that when we use 1 CPU, mutex lock runs much faster than spinlock because instead of spinning, mutex lock will release the CPU to another thread.

To implement all these user defined system calls, we have made the added the following pieces of codes to the respective files -

Usys.s

```
32 SYSCALL(thread_create)
33 SYSCALL(thread_join)
34 SYSCALL(thread_exit)
```

User.h

```
27 int thread_create(void (*)(void*), void*, void*);
28 int thread_join(void);
29 int thread_exit(void);
```

Syscall.h

```
23 #define SYS_thread_create 22
24 #define SYS_thread_exit 23
25 #define SYS_thread_join 24
```

Syscall.c

```
133 [SYS_thread_create]    sys_thread_create,  
134 [SYS_thread_join]     sys_thread_join,  
135 [SYS_thread_exit]     sys_thread_exit,
```

```
106 extern int sys_thread_create(void);  
107 extern int sys_thread_join(void);  
108 extern int sys_thread_exit(void);
```

Defs.h

```
124 //thread  
125 int thread_create(void (*)(void*),void*,void*);  
126 int thread_join(void);  
127 int thread_exit(void);  
128
```

Makefile

```
168 UPROGS=\n169     _cat\  
170     _echo\  
171     _forktest\  
172     _grep\  
173     _init\  
174     _kill\  
175     _ln\  
176     _ls\  
177     _mkdir\  
178     _rm\  
179     _sh\  
180     _stressfs\  
181     _wc\  
182     _zombie\  
183     _thread\  
184     _Drawtest\  
185
```